Dhruv Bhavsar

+91 9558452967 | dbhavsar9898@gmail.com | linkedin.com/in/dhruv4real | github.com/dhruvbhavsar0612 | India

Summary

AI Engineer and builder specialising in production agentic pipelines, real-time voice AI, and event-driven ML infrastructure. 2+ years architecting and shipping AI products across healthcare, enterprise automation, and telephony; currently leading AI engineering at Wappnet Systems with a patent-pending routing system and 500+ daily live voice agent calls in production.

Education

Master of Engineering, Artificial Intelligence — Sep 2026 (Incoming)
University of California, Los Angeles (UCLA) — *Los Angeles, CA*

Bachelor of Engineering, Computer Science (AI/ML) — 2024
Gujarat Technological University — *Ahmedabad, India*
- CGPA: 9.6/10

Experience

SDE 2, AI & Data Science — Jan 2025 – Present
Wappnet Systems — *Ahmedabad, India*
- Lead a team of 6–7 AI engineers: plan system architecture from client requirements, conduct code reviews, manage client communication across the full delivery cycle, and drive adoption of AI-assisted development (Claude Code, Codex, GitHub Copilot) across the department
- Deployed voice AI agents for pharmacy clients handling **500+ inbound calls daily**: answer general queries, check refill availability, escalate to staff, and act as real-time multilingual mediators on contact-centre calls; reduced reception call volume by **70%**. Developed admin platform for observability and configuration in Next.
- Developed an AI frontdesk: businesses connect documents, websites, or Google Places data to a deployable chat widget resolving repetitive queries at point of contact; onboarded **20+ customers** on a live voice-enabled plan
- Built maritime routing API for the Finnish government: Contraction Hierarchies on a 50,000+ node graph cut query latency from **2s to under 100ms (p99)**; MVP accepted for **patent filing under Finnish regulations**

AI Engineer (Employee of the Year 2024, 4× Quarterly Awards) — July 2024 – Dec 2024
Wappnet Systems — *Ahmedabad, India*
- Architected **osora.ai** enterprise agentic copilot on AWS Bedrock: multi-agent state machine with Redis for active memory and PostgreSQL for long-term recall, reducing token consumption by **63%** via separated Reasoning and IO agents, with custom MCP servers for Slack, Jira, Notion, Google Workspace, and Outlook
- Delivered **nolea.ai** healthcare CRM copilot: fine-tuned embedding models with GPU-accelerated collaborative filtering on Elasticsearch, plus a RAG assistant for natural-language queries across the healthcare database

Data Engineering Intern — Jan 2024 – Jun 2024
Wappnet Systems — *Ahmedabad, India*
- Engineered 53 Apache Airflow DAGs to ingest and transform multi-source healthcare records into production-ready profiles; implemented DBSCAN geographic clustering for daily recommendations, improving match accuracy by **35%**
- Built the team's first RAG system using Approximate Nearest Neighbor search over people profiles; deployed auto-scaling geocoding service on AWS Lambda using Polars and PySpark

Research Intern — Oct 2023 – Jan 2024
Indian Space Research Organisation (ISRO) — *Ahmedabad, India*
- Trained an LSTM+CNN hybrid model for solar insolation prediction fusing 23 satellite sensor instruments, outperforming ARIMA/SARIMA baselines by **40%** during high-fluctuation weather events
- Engineered the full ETL pipeline over 200GB+ of satellite raster archives and deployed a continuous retraining pipeline with NVIDIA CuDNN for zero-intervention production updates

Projects

TeleAI Indic (teleai.tech) 2025
Building side venture *Python, FastAPI, OpenAI Realtime, Twilio, React*

- Asymmetric voice AI platform: humans speak, AI responds in real-time with concise written output only. Built on the HCI principle that speech is faster than typing but reading is faster than listening; live at voice.teleai.tech and chat.teleai.tech with AWS deployment and GitHub Actions CI/CD

fastapi-smith – PyPI – GitHub 2025
Open Source *Python, PyPI, GitHub Actions*

- Production-ready FastAPI project generator with automated versioned releases to PyPI.

rustlette 2025
Open Source *Rust*

- Rust reimplementation of the Starlette ASGI framework achieving a **10% improvement** in memory and throughput over the Python baseline via zero-cost async abstractions

Technical Skills

Languages: Python, Rust, TypeScript, JavaScript, SQL
AI & ML: Agentic AI, MCP (Model Context Protocol), OpenAI Agents SDK, AWS Bedrock, RAG Systems, LLM Fine-tuning, Embedding Models, Vector Databases, Voice AI (Twilio, OpenAI Realtime), LSTM+CNN, TensorFlow, PyTorch, Keras, Scikit-learn, Elasticsearch, NVIDIA CuDNN
Backend & Infrastructure: FastAPI, PostgreSQL, Redis, MongoDB, Docker, AWS (EC2, Lambda, Bedrock), GitHub Actions, SQLAlchemy, Alembic, Microservices, Apache Airflow
Data Engineering: ETL Pipelines, PySpark, Polars, GeoPandas, NetworkX, pandas, NumPy, OpenCV2